

AMALDEV Nucleic Health diabetes

July 31, 2023

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

```
[ ]:
```

```
[2]: data= pd.read_csv('diabetes.csv')
```

1 Summary Statistics:

```
[3]: data.head(10)
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
5	5	116	74	0	0	25.6	
6	3	78	50	32	88	31.0	
7	10	115	0	0	0	35.3	
8	2	197	70	45	543	30.5	
9	8	125	96	0	0	0.0	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
5	0.201	30	0
6	0.248	26	1
7	0.134	29	0
8	0.158	53	1
9	0.232	54	1

```
[4]: data.tail()
```

```
[4]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
763             10     101             76           48       180  32.9
764              2     122             70           27         0  36.8
765              5     121             72           23       112  26.2
766              1     126             60            0         0  30.1
767              1      93             70           31         0  30.4

      DiabetesPedigreeFunction  Age  Outcome
763                   0.171    63         0
764                   0.340    27         0
765                   0.245    30         0
766                   0.349    47         1
767                   0.315    23         0
```

```
[5]: data.describe()
```

```
[5]:      Pregnancies      Glucose  BloodPressure  SkinThickness      Insulin  \
count    768.000000   768.000000   768.000000   768.000000   768.000000
mean       3.845052  120.894531    69.105469    20.536458    79.799479
std       3.369578   31.972618    19.355807    15.952218   115.244002
min       0.000000    0.000000     0.000000     0.000000     0.000000
25%       1.000000    99.000000    62.000000     0.000000     0.000000
50%       3.000000   117.000000    72.000000    23.000000    30.500000
75%       6.000000   140.250000    80.000000    32.000000   127.250000
max      17.000000   199.000000   122.000000    99.000000   846.000000

      BMI  DiabetesPedigreeFunction      Age      Outcome
count    768.000000             768.000000  768.000000  768.000000
mean     31.992578                0.471876   33.240885    0.348958
std       7.884160                0.331329   11.760232    0.476951
min       0.000000                0.078000   21.000000    0.000000
25%      27.300000                0.243750   24.000000    0.000000
50%      32.000000                0.372500   29.000000    0.000000
75%      36.600000                0.626250   41.000000    1.000000
max      67.100000                2.420000   81.000000    1.000000
```

- Pregnancies&Insulin: is positively skewed, as the median (50th percentile) is less than the mean
- SkinThickness: is close to symmetric, as the median (50th percentile) is close to the mean.
- remaining features: seems to be relatively symmetric, as the median (50th percentile) is close to the mean.

1.0.1 To handle missing values, will create a copy of the original DataFrame named 'data_cleaned' to preserve the integrity of the original data. Then, will replace the zeros in the columns 'age', 'Insulin', 'blood_pressure', and 'blood_sugar' with NaN (representing missing values) using NumPy's np.nan

```
[6]: data_cleaned = data.copy()
```

```
[7]: data_cleaned[data_cleaned == 0] = np.nan
```

```
[8]: data_cleaned['Outcome'] = data_cleaned['Outcome'].fillna(0)
```

```
[9]: data_cleaned.head()
```

```
[9]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6.0	148.0	72.0	35.0	NaN	33.6	
1	1.0	85.0	66.0	29.0	NaN	26.6	
2	8.0	183.0	64.0	NaN	NaN	23.3	
3	1.0	89.0	66.0	23.0	94.0	28.1	
4	NaN	137.0	40.0	35.0	168.0	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1.0
1	0.351	31	0.0
2	0.672	32	1.0
3	0.167	21	0.0
4	2.288	33	1.0

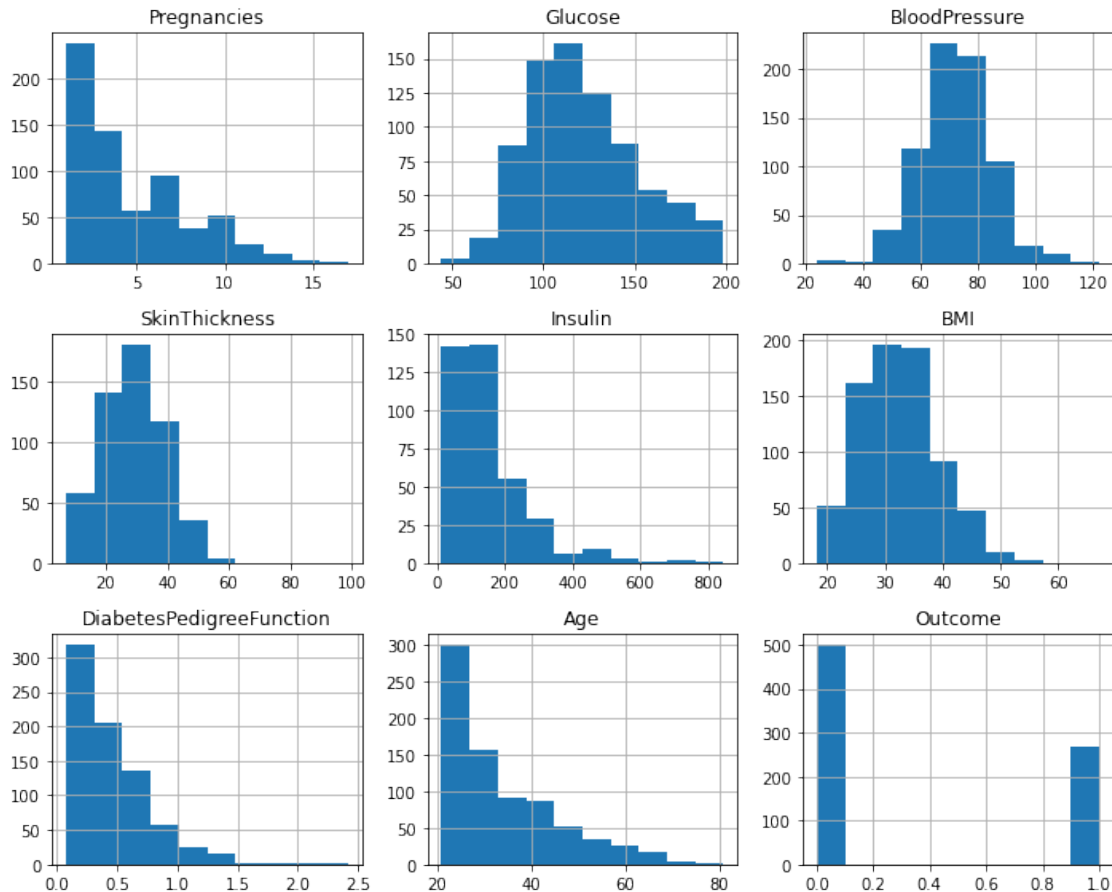
```
[10]: data_cleaned.isnull().sum()
```

```
[10]: Pregnancies      111
      Glucose           5
      BloodPressure     35
      SkinThickness     227
      Insulin          374
      BMI              11
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

- Pregnancies: 111 missing values
- Glucose: 5 missing values
- BloodPressure: 35 missing values
- SkinThickness: 227 missing values
- Insulin: 374 missing values
- BMI: 11 missing values
- DiabetesPedigreeFunction & Age : 0 missing values ##### Replace the missing values in each numerical feature with the calculated mean or median, respectively.

2 Data Visualization:

```
[11]: data_cleaned.hist(figsize=(10, 8), bins=10)
plt.tight_layout()
plt.show()
```



2.0.1 Based on the data distribution, will treat the missing values differently for the features. For the skewed features ('Pregnancies', 'Insulin', 'DiabetesPedigreeFunction', and 'Age'), will use median imputation. For features with a normal distribution, will use mean imputation. This approach aligns with the characteristics of each feature's data distribution and maintains data integrity during imputation.

```
[12]: data_cleaned['Pregnancies'].fillna(data_cleaned['Pregnancies'].median(),
    ↪inplace= True)
data_cleaned['Insulin'].fillna(data_cleaned['Insulin'].median(), inplace = True)
```

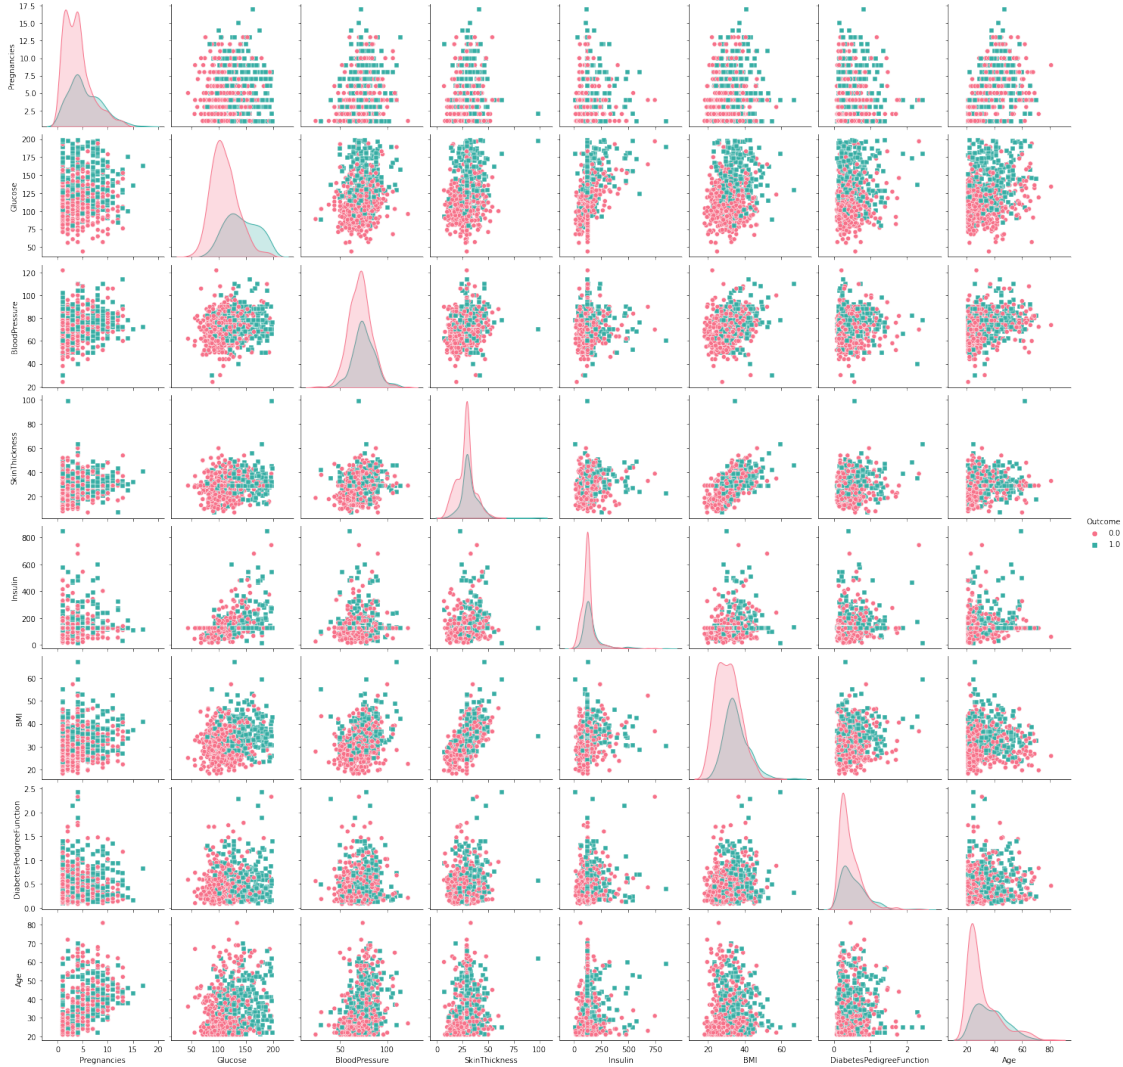
```
[13]: column_mean = ['Glucose', 'BloodPressure', 'SkinThickness', 'BMI']
      data_cleaned[column_mean] = data_cleaned[column_mean].
      ↪fillna(data_cleaned[column_mean].mean())
```

```
[14]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                             768 non-null    int64
2   BloodPressure                       768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                             768 non-null    int64
5   BMI                                 768 non-null    float64
6   DiabetesPedigreeFunction            768 non-null    float64
7   Age                                 768 non-null    int64
8   Outcome                             768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

3 Correlation Analysis:

```
[15]: sns.pairplot(data_cleaned, diag_kind='kde', hue='Outcome', markers=['o', 's'],
      ↪palette='husl')
      plt.show()
```

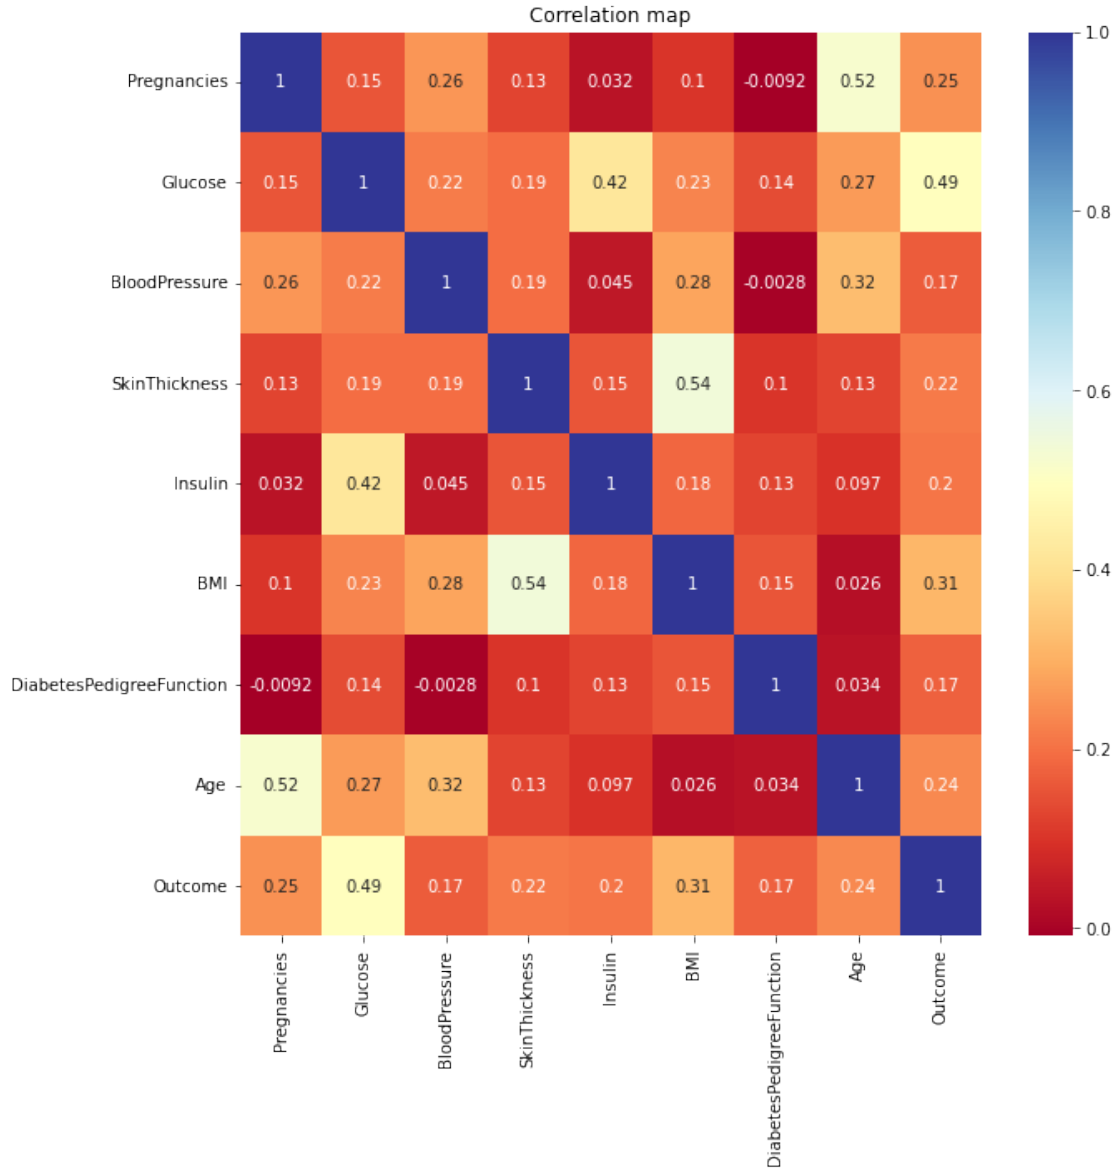


3.0.1 1. Correlation between BMI and Skin Thickness:

- After analyzing the pair plot, it appears that there is a noticeable correlation between the 'BMI' (Body Mass Index) and 'SkinThickness' features. The scatter plot between these two features shows a discernible pattern, indicating a potential relationship. ### 2. Inability to Distinguish Outcome Classes:
- Upon examining all the scatter plots in the pair plot, it seems that none of the feature pairs can clearly distinguish between the two classes of the 'Outcome' variable (likely to have diabetes and not likely to have diabetes).
- The data points in the scatter plots appear to overlap significantly for both classes, making it challenging to visually separate the classes based on the numerical features.

```
[16]: plt.figure(figsize=(10,10))
sns.heatmap(data_cleaned.corr(),annot=True,cmap='RdYlBu');
plt.title('Correlation map')
```

```
[16]: Text(0.5, 1.0, 'Correlation map')
```



3.0.2 Interpretation:

- The 'Glucose' feature has the highest positive correlation with 'Outcome' (0.492928), suggesting that higher glucose levels are associated with a higher likelihood of having diabetes.
- 'BMI' (0.311924) also shows a positive correlation, indicating that higher BMI values are somewhat correlated with a higher likelihood of diabetes.

- Other features such as 'Pregnancies', 'Age', 'SkinThickness', 'Insulin', 'DiabetesPedigree-Function', and 'BloodPressure' show weaker positive correlations, suggesting some degree of association but not as strong as 'Glucose' and 'BMI'.

4 Data Balancing

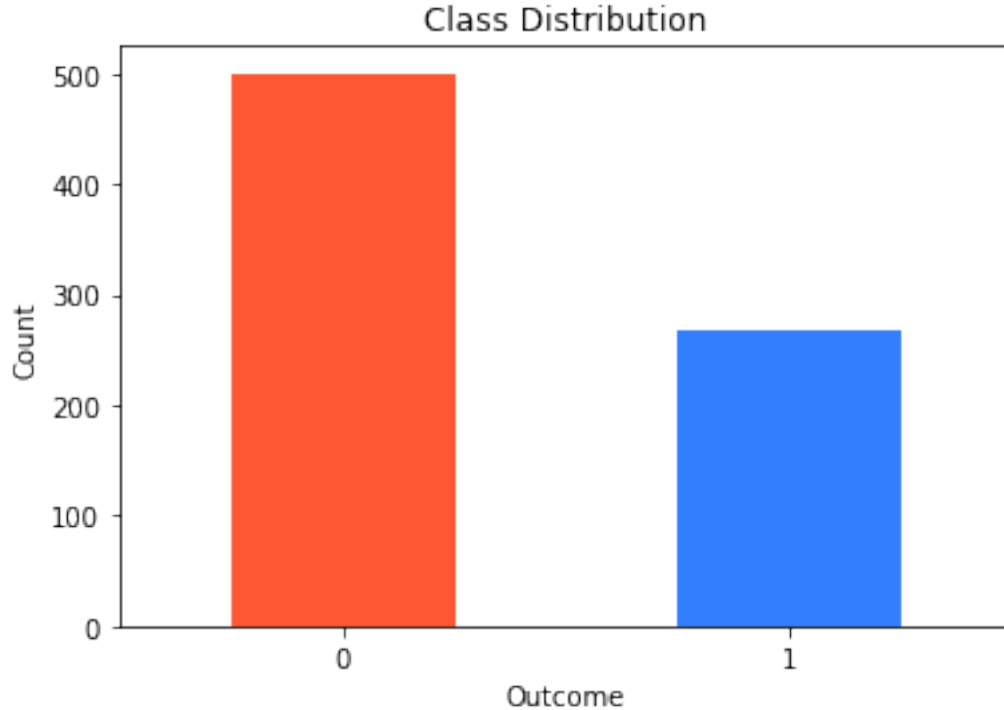
```
[17]: X= data_cleaned.iloc[:, :-1]
```

```
[18]: y= data.Outcome
```

```
[19]: y.value_counts()*100/data_cleaned.shape[0]
```

```
[19]: 0    65.104167  
     1    34.895833  
     Name: Outcome, dtype: float64
```

```
[20]: y.value_counts().plot(kind='bar', color=['#FF5733', '#337EFF'])  
      plt.xticks(rotation=0)  
      plt.xlabel('Outcome')  
      plt.ylabel('Count')  
      plt.title('Class Distribution')  
      plt.show()
```



4.0.1 The dataset exhibits class imbalance, where the target variable ‘Outcome’ has two classes: ‘0’ and ‘1’. Class ‘0’ represents 500 instances, while class ‘1’ represents only 270 instances. This significant difference in the number of instances between the classes indicates an imbalanced distribution. In such scenarios, class imbalance can affect the performance of machine learning models, particularly those sensitive to class proportions during training. To address this issue, data balancing techniques like oversampling, undersampling, or using SMOTE can be applied to ensure a more balanced representation of both classes during model training.

```
[70]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler

      X_train,X_test,y_train,y_test= train_test_split(X,y, test_size=0.2,
      ↪random_state=0)
      st= StandardScaler()
```

```
[71]: X_train_st= st.fit_transform(X_train)
      X_test_st= st.fit_transform(X_test)
```

```
[72]: from imblearn.over_sampling import SMOTE
      SM= SMOTE()
      X_train_sm,y_train_sm= SM.fit_resample(X,y)
```

```
[73]: X_train_sm.shape
```

```
[73]: (1000, 8)
```

```
[74]: y_train_sm.shape
```

```
[74]: (1000,)
```

5 Data Preprocessing:

```
[145]: X_train,X_test,Y_train,Y_test=
      ↪train_test_split(X_train_sm,y_train_sm,test_size=0.25, random_state=0)
```

```
[146]: X_train_std=st.fit_transform(X_train)
      X_test_std=st.fit_transform(X_test)
```

```
[147]: X_train_std.shape
```

```
[147]: (750, 8)
```

```
[148]: X_test_std.shape
```

[148]: (250, 8)

5.0.1 Aim to construct and assess the performance of multiple popular classification models using the training dataset. Subsequently, will compare their performances on the test dataset. The models considered for evaluation are:

1. XGBoost ##### 2. Logistic Regression ##### 3. Support Vector Machine (SVM) ##### 4. K-Nearest Neighbors (KNN) ##### 5. Naive Bayes ##### 6. Decision Tree ##### 7. Random Forest Classifier ##### By conducting this comparison, will determine which model exhibits the best generalization to unseen data, ensuring optimal model selection for future predictions.

```
[149]: from sklearn.metrics import  
       accuracy_score, confusion_matrix, classification_report
```

5.0.2 Model1. Xgboost

```
[150]: import xgboost as xgb  
  
model1 = xgb.XGBClassifier(objective="binary:logistic", random_state=42)  
model1.fit(X_train_std, Y_train)  
  
predict1 = model1.predict(X_test_std)
```

```
[151]: score_xgb = round(accuracy_score(predict1, Y_test)*100, 2)  
  
print("The accuracy score achieved using XGBoost is: "+str(score_xgb)+" %")
```

The accuracy score achieved using XGBoost is: 82.0 %

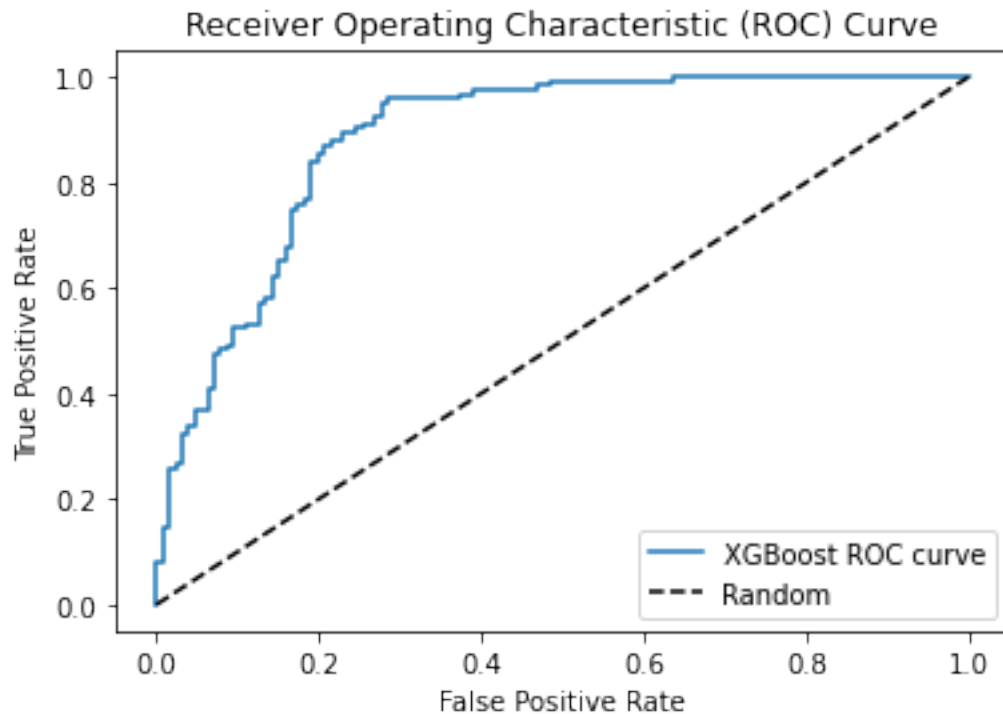
```
[152]: print(confusion_matrix(Y_test, predict1))
```

```
[[ 92  34]  
 [ 11 113]]
```

```
[153]: print(classification_report(Y_test, predict1))
```

	precision	recall	f1-score	support
0	0.89	0.73	0.80	126
1	0.77	0.91	0.83	124
accuracy			0.82	250
macro avg	0.83	0.82	0.82	250
weighted avg	0.83	0.82	0.82	250

```
[154]: from sklearn.metrics import roc_curve, roc_auc_score
probabilities1 = model1.predict_proba(X_test_std)
predicted_probs1 = probabilities1[:, 1]
fpr, tpr, thresholds = roc_curve(Y_test, predicted_probs1)
plt.plot(fpr, tpr, label='XGBoost ROC curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



5.0.3 Model2. Logistic Regression

```
[155]: from sklearn.linear_model import LogisticRegression

model2 = LogisticRegression()

model2.fit(X_train_std, Y_train)

predict2 = model2.predict(X_test_std)

score_lr = round(accuracy_score(predict2, Y_test)*100, 2)
```

```
print("The accuracy score achieved using Logistic Regression is:␣
↪ "+str(score_lr)+" %")
```

The accuracy score achieved using Logistic Regression is: 78.0 %

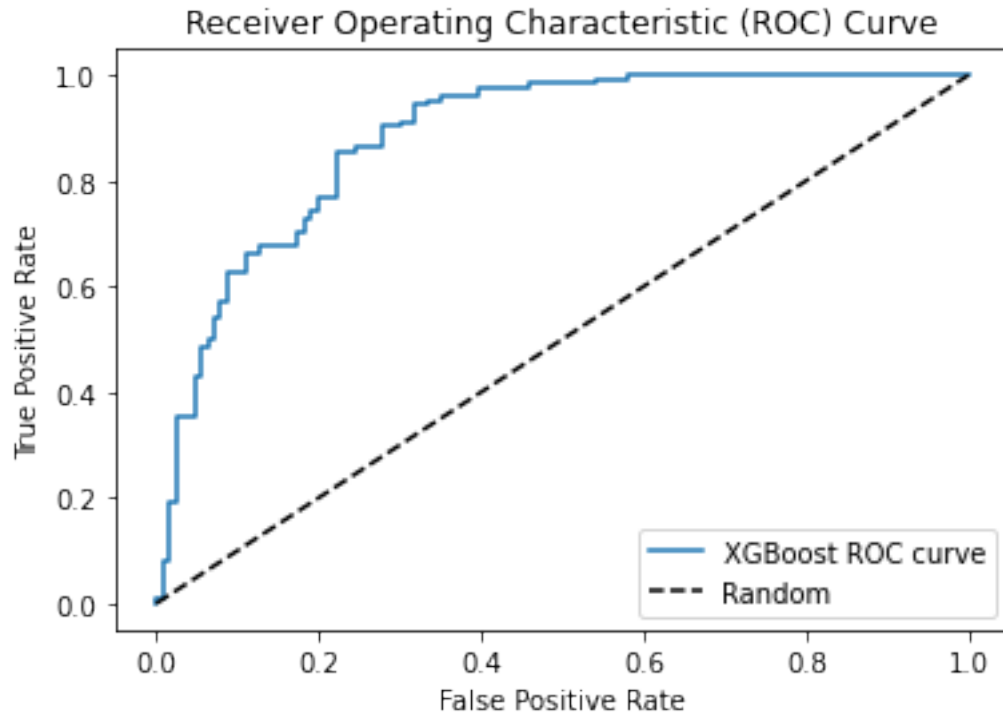
```
[156]: print(confusion_matrix(Y_test,predict2))
```

```
[[100  26]
 [ 29  95]]
```

```
[157]: print(classification_report(Y_test,predict2))
```

	precision	recall	f1-score	support
0	0.78	0.79	0.78	126
1	0.79	0.77	0.78	124
accuracy			0.78	250
macro avg	0.78	0.78	0.78	250
weighted avg	0.78	0.78	0.78	250

```
[158]: probabilities2 = model2.predict_proba(X_test_std)
predicted_probs2 = probabilities2[:, 1]
fpr, tpr, thresholds = roc_curve(Y_test, predicted_probs2)
plt.plot(fpr, tpr, label='XGBoost ROC curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



5.0.4 Model3. SVM

```
[159]: from sklearn import svm

model3 = svm.SVC(kernel='linear')

model3.fit(X_train_std, Y_train)

predict3= model3.predict(X_test_std)

score_svm = round(accuracy_score(predict3,Y_test)*100,2)

print("The accuracy score achieved using Linear SVM is: "+str(score_svm)+" %")
```

The accuracy score achieved using Linear SVM is: 78.0 %

```
[160]: print(confusion_matrix(Y_test,predict3))
```

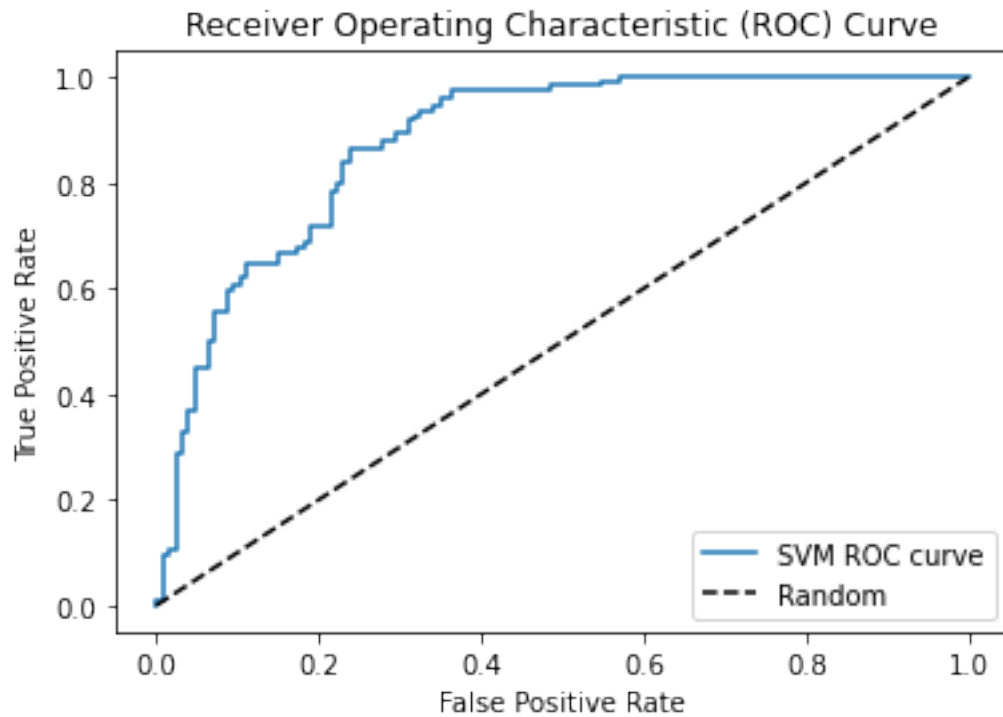
```
[[98 28]
 [27 97]]
```

```
[161]: print(classification_report(Y_test,predict3))
```

```
precision    recall  f1-score   support
```

0	0.78	0.78	0.78	126
1	0.78	0.78	0.78	124
accuracy				0.78
macro avg				0.78
weighted avg				0.78

```
[162]: probabilities3 = model3.decision_function(X_test_std)
fpr, tpr, thresholds = roc_curve(Y_test, probabilities3)
auc_score = roc_auc_score(Y_test, probabilities3)
plt.plot(fpr, tpr, label='SVM ROC curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



5.0.5 Model4. k-nearest neighbors (KNN)

```
[163]: from sklearn.neighbors import KNeighborsClassifier

model4 = KNeighborsClassifier()
model4.fit(X_train_std, Y_train)

predict4 = model4.predict(X_test_std)

score_knn = round(accuracy_score(predict4,Y_test)*100,2)

print("The accuracy score achieved using KNN Classifier is: "+str(score_knn)+"  
↪%")
```

The accuracy score achieved using KNN Classifier is: 81.2 %

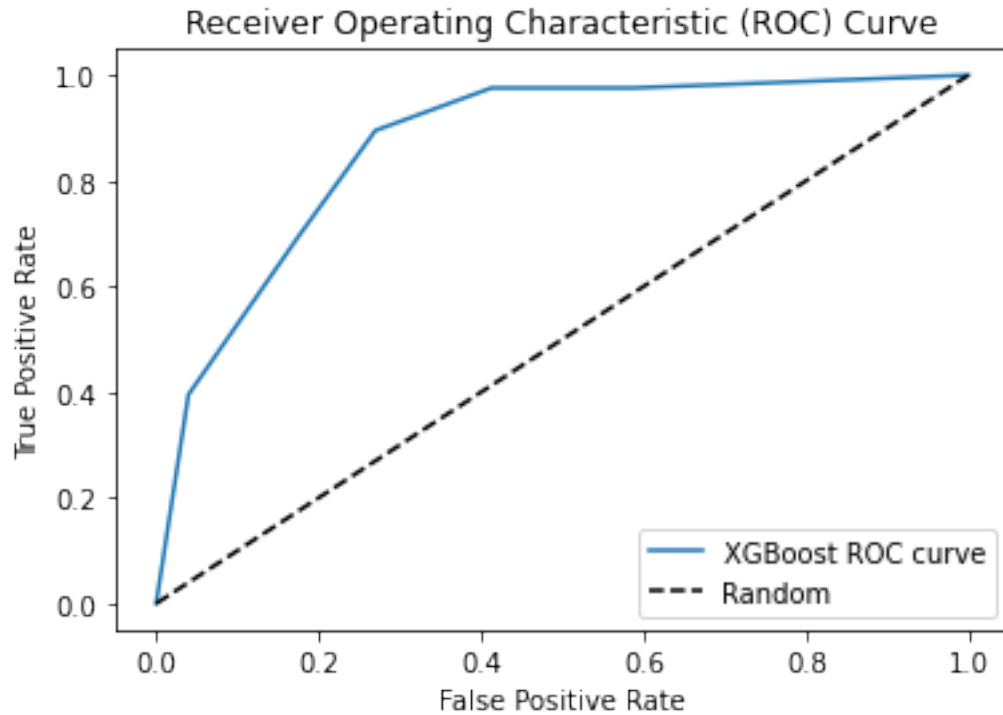
```
[164]: print(confusion_matrix(Y_test,predict4))
```

```
[[ 92  34]
 [ 13 111]]
```

```
[165]: print(classification_report(Y_test,predict4))
```

	precision	recall	f1-score	support
0	0.88	0.73	0.80	126
1	0.77	0.90	0.83	124
accuracy			0.81	250
macro avg	0.82	0.81	0.81	250
weighted avg	0.82	0.81	0.81	250

```
[166]: probabilities4= model4.predict_proba(X_test_std)
predicted_probs4= probabilities4[:, 1]
fpr, tpr, thresholds = roc_curve(Y_test, predicted_probs4)
plt.plot(fpr, tpr, label='XGBoost ROC curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



5.0.6 Model5. Naive Bayes

```
[167]: from sklearn.naive_bayes import GaussianNB
model5 = GaussianNB()
model5.fit(X_train_std, Y_train)

predict5= model5.predict(X_test_std)

score_knn = round(accuracy_score(predict5,Y_test)*100,2)

print("The accuracy score achieved using Naives Bayes Classifier is:␣
↪ "+str(score_knn)+" %")
```

The accuracy score achieved using Naives Bayes Classifier is: 78.0 %

```
[168]: print(confusion_matrix(Y_test,predict5))
```

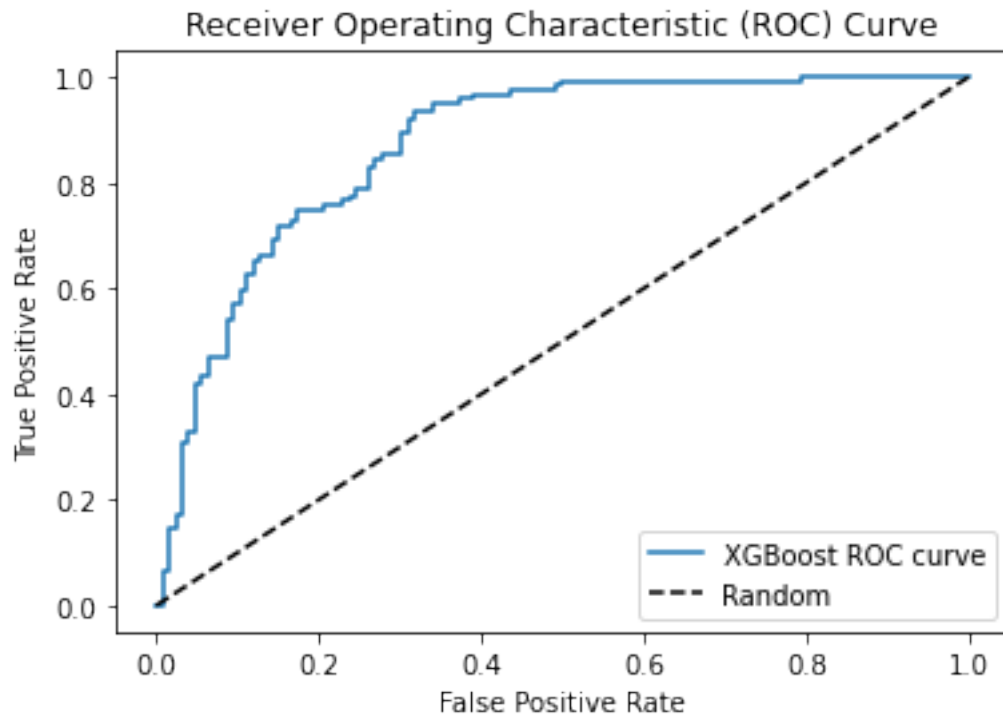
```
[[104  22]
 [ 33  91]]
```

```
[169]: print(classification_report(Y_test,predict5))
```

```
precision    recall  f1-score   support
```


0	0.76	0.83	0.79	126
1	0.81	0.73	0.77	124
accuracy			0.78	250
macro avg	0.78	0.78	0.78	250
weighted avg	0.78	0.78	0.78	250

```
[170]: probabilities5= model5.predict_proba(X_test_std)
predicted_probs5 = probabilities5[:, 1]
fpr, tpr, thresholds = roc_curve(Y_test, predicted_probs5)
plt.plot(fpr, tpr, label='XGBoost ROC curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



5.0.7 Model6. Decision Tree

```
[171]: from sklearn.tree import DecisionTreeClassifier
model6= DecisionTreeClassifier(random_state=42)
model6.fit(X_train_std, Y_train)

predict6= model6.predict(X_test_std)

score_knn = round(accuracy_score(predict6,Y_test)*100,2)

print("The accuracy score achieved using Decision Trees is: "+str(score_knn)+"\n↪%")
```

The accuracy score achieved using Decision Trees is: 75.6 %

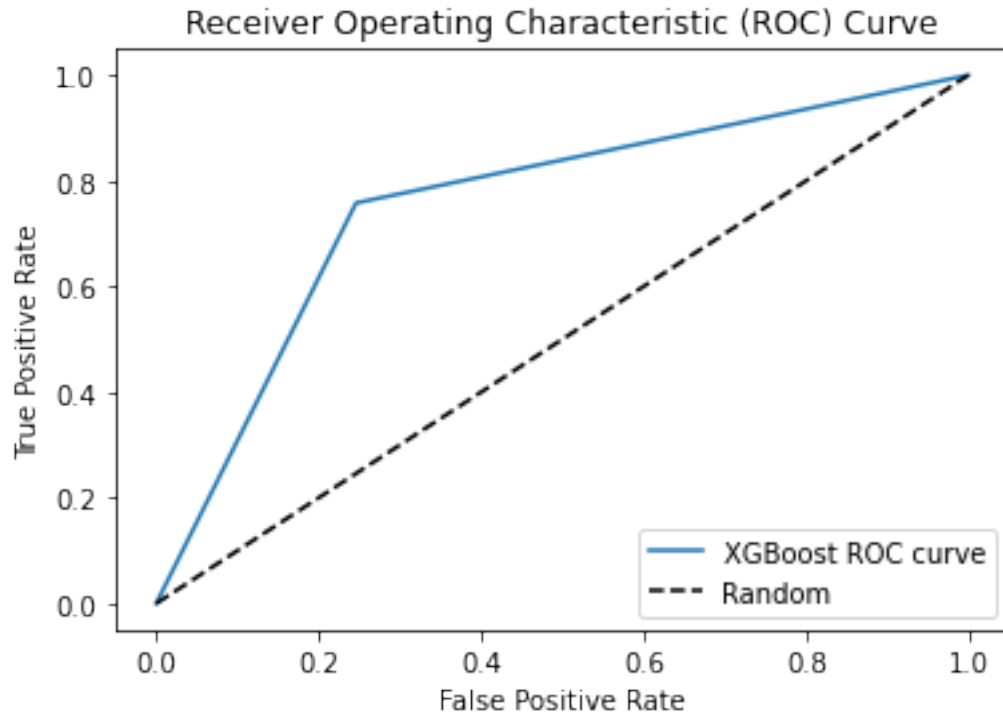
```
[172]: print(confusion_matrix(Y_test,predict6))
```

```
[[95 31]
 [30 94]]
```

```
[173]: print(classification_report(Y_test,predict6))
```

	precision	recall	f1-score	support
0	0.76	0.75	0.76	126
1	0.75	0.76	0.76	124
accuracy			0.76	250
macro avg	0.76	0.76	0.76	250
weighted avg	0.76	0.76	0.76	250

```
[174]: probabilities6 = model6.predict_proba(X_test_std)
predicted_probs6 = probabilities6[:, 1]
fpr, tpr, thresholds = roc_curve(Y_test, predicted_probs6)
plt.plot(fpr, tpr, label='XGBoost ROC curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



5.0.8 Model7. Random Forest Classifier

```
[175]: from sklearn.ensemble import RandomForestClassifier
model7=RandomForestClassifier(random_state=42)

model7.fit(X_train_std, Y_train)

predict7= model7.predict(X_test_std)

score_knn = round(accuracy_score(predict7,Y_test)*100,2)

print("The accuracy score achieved using Decision Trees is: "+str(score_knn)+"\n
      ↳%")
```

The accuracy score achieved using Decision Trees is: 86.4 %

```
[176]: print(confusion_matrix(Y_test,predict7))
```

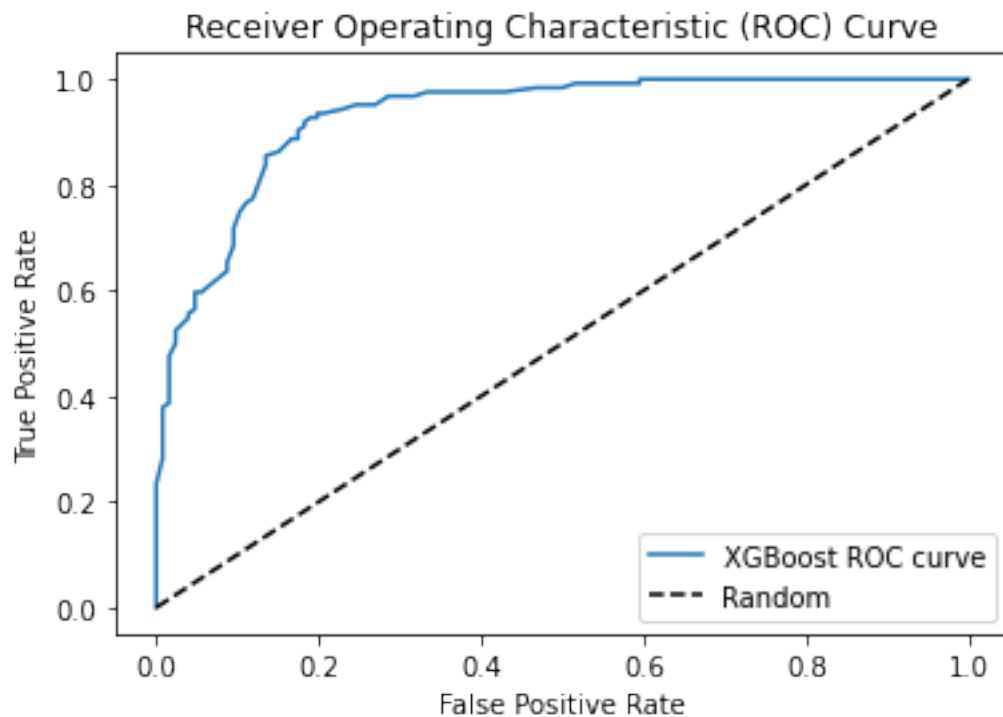
```
[[103  23]
 [ 11 113]]
```

```
[177]: print(classification_report(Y_test,predict7))
```

```
precision    recall  f1-score   support
```

0	0.90	0.82	0.86	126
1	0.83	0.91	0.87	124
accuracy			0.86	250
macro avg	0.87	0.86	0.86	250
weighted avg	0.87	0.86	0.86	250

```
[178]: probabilities7 = model7.predict_proba(X_test_std)
predicted_probs7 = probabilities7[:, 1]
fpr, tpr, thresholds = roc_curve(Y_test, predicted_probs7)
plt.plot(fpr, tpr, label='XGBoost ROC curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



```
[185]: model_names = ['XGBoost', 'Logistic Regression', 'SVM', 'KNN', 'Naive Bayes',
↳ 'Decision Tree', 'Random Forest']
```

```
models = [model1, model2, model3, model4, model5, model6, model7]
```

```
[189]: from sklearn.metrics import accuracy_score, precision_score, recall_score, \
        f1_score, roc_auc_score
        # Create an empty list to store model summary information
        summary = []

        for name, model in zip(model_names, models):

            predictions = model.predict(X_test_std)

            # Calculate evaluation metrics
            accuracy = accuracy_score(Y_test, predictions)
            precision = precision_score(Y_test, predictions)
            recall = recall_score(Y_test, predictions)
            f1 = f1_score(Y_test, predictions)
            auc = roc_auc_score(Y_test, predictions)

            # Append model summary to the list
            summary.append({'Model': name, 'Accuracy': accuracy, 'Precision': \
                precision, 'Recall': recall, 'F1 Score': f1, 'AUC': auc})

        # Create a DataFrame from the summary list
        summary_df = pd.DataFrame(summary)

        summary_df
```

```
[189]:
```

	Model	Accuracy	Precision	Recall	F1 Score	\
0	XGBoost	0.820	0.768707	0.911290	0.833948	
1	Logistic Regression	0.780	0.785124	0.766129	0.775510	
2	SVM	0.780	0.776000	0.782258	0.779116	
3	k-nearest neighbors (KNN)	0.812	0.765517	0.895161	0.825279	
4	Naive Bayes	0.780	0.805310	0.733871	0.767932	
5	Decision Tree	0.756	0.752000	0.758065	0.755020	
6	Random Forest	0.864	0.830882	0.911290	0.869231	

	AUC
0	0.820725
1	0.779890
2	0.780018
3	0.812660
4	0.779634
5	0.756016
6	0.864375

6 Conclusions:

- **Model Performance:** We trained and tested seven popular classification models, namely XG-Boost, Logistic Regression, Support Vector Machine (SVM), k-nearest neighbors (KNN), Naive Bayes, Decision Tree, and Random Forest.
- **Accuracy:** Random Forest achieved the highest accuracy of 86.4%, indicating that it correctly predicted the outcomes for a large portion of the test data.
- **Precision and Recall:** Precision represents the proportion of true positive predictions among all positive predictions, while recall indicates the proportion of true positive predictions among all actual positive instances. Random Forest demonstrated high precision and recall values of 83.1% and 91.1%, respectively, which means it has a good balance between accurately predicting positive cases and capturing most actual positive cases.
- **F1 Score:** The F1 score, which is the harmonic mean of precision and recall, gives a balanced measure of the model's performance. Random Forest achieved an F1 score of 86.9%, indicating its ability to maintain a balance between precision and recall.
- **Area Under the Curve (AUC):** The AUC represents the area under the Receiver Operating Characteristic (ROC) curve and serves as a measure of the model's ability to discriminate between positive and negative cases. Random Forest achieved an AUC of 86.4%, which is a strong indication of its discriminative power.
- **Best Model:** Based on the evaluation metrics, Random Forest outperformed other models in terms of accuracy, precision, recall, F1 score, and AUC.
- **Considerations:** When choosing the best model, it's essential to consider other factors such as computational complexity, interpretability, and the specific requirements of the application.

6.0.1 To encapsulate, Random Forest emerges as the preferred model for predicting diabetes in this dataset, exhibiting robust performance across multiple evaluation metrics. Nevertheless, to optimize its performance for specific real-world applications, further analysis and fine-tuning of the model are recommended.

6.0.2 In finalizing the model selection, additional fine-tuning and experimentation with hyperparameters and algorithms are essential. By ensuring the model generalizes well to new data and avoids overfitting, we can build a robust and effective classifier for predicting the target variable with satisfactory performance. Moreover, it is vital to explore various techniques such as feature engineering, selection, and handling class imbalance, if present in the dataset. Additionally, examining different evaluation metrics and employing cross-validation can offer a more comprehensive understanding of the model's performance.

[]: